

AI & RAG

Unistack Solutions develops Retrieval-Augmented Generation (RAG) systems using LangChain, LangGraph, OpenAI-compatible models, embeddings and vector databases such as ChromaDB, Qdrant, FAISS and Pinecone.

Typical workflow: Load Documents → Chunk Text → Create Embeddings → Store in Vector DB → Similarity Search → Generate LLM Response.